

Excavation Field Guide

Part One

Flat dialogue is the surface symptom of a substrate problem. The character says exactly what they mean because the writer has not yet identified what the character cannot say — and what they cannot say is always a function of what they are protecting.

THE RESEARCH

"People rarely say exactly what they mean. They hint. They deflect. They talk around the thing they actually want to say." — Savannah Gilbo, writing instructor and developmental editor, on the mechanics of subtext in fiction dialogue.

This is not a stylistic preference. It is how human beings operate under pressure. Research in behavioral psychology and linguistics consistently shows that self-disclosure is regulated — people reveal only what their current emotional state permits them to reveal. They don't withhold because they're interesting. They withhold because exposure is dangerous.

When you write a character who says exactly what they feel, you have inadvertently removed the protective instinct that makes human speech authentic. The reader doesn't consciously notice the mechanism. But something fails to land. The scene feels constructed. The character feels performed.

This means that craft fixes — applied to the surface — leave the underlying problem unchanged. The line gets technically cleaner. The protection is still missing. The character still sounds like someone who has no reason to hide.

THE PRACTICAL IMPLICATION

Before you rewrite a flat line, name what the character is protecting. If you cannot name it, any rewrite will be cosmetic. The substrate discovery comes first. The language follows.

Part Two

Socrates never told anyone the answer. His method was to ask a question, receive an answer, and then ask the question underneath that answer — until the person's own knowledge became visible to them. This is not a technique. It is a premise about where truth lives.

Your AI coach operates on the same premise. Understanding the three mechanics of a Socratic session allows you to recognize when each is working, and to intervene when it isn't.

01 Questions, not answers

The AI's role in a coaching session is to ask, never to provide. Every time it generates a rewrite, a suggestion, or a characterization, it has left the Socratic role. The moment you receive content instead of a question, the session has shifted from excavation to fabrication. You can restore it with a single override line.

02 The first answer is surface

Cognitive science research on insight and problem-solving (Dijksterhuis & Meurs, 2006; Bowden & Jung-Beeman, 2003) consistently shows that unconscious processing produces different — often richer — answers than conscious deliberation. Your first answer to a coaching question is produced by conscious deliberation. It is usually correct. It is almost never the substrate. Plan to reach a third or fourth answer before you trust that you have found the real thing.

03 Discomfort is a signal, not a problem

The substrate answer tends to arrive with a quality of mild recognition — something you knew but hadn't named. It often also arrives with slight resistance. This is neurologically consistent with the experience of retrieving implicit memories: the material is true, and the body registers it before the mind fully catches up. If an answer produces no friction at all, you are likely still in surface territory. If it produces significant distress, you may have overshot. The useful signal is a quiet, specific discomfort — something exact that you were not planning to write.

Part Three

Two failure modes account for nearly every unproductive Socratic session. Recognizing them in the moment is the skill. The recovery is always a single line.

FAILURE MODE	RECOVERY
The AI drifts into content generation — it rewrites the line, suggests what the character "really means," or begins producing options.	"Do not rewrite the line. Ask me what my character is afraid to say."
You accept a surface answer and the AI moves on without pushing deeper.	"Stay on this question. Ask me again from a different angle."
Your answers are one sentence and the session feels thin.	Force yourself to write three sentences. Your third will be truer than your first.
The session produces clever-sounding answers that have no felt weight.	Ask: "What would this character never say in this scene?" The inverse question often reaches the substrate faster.
You feel stuck and aren't sure why the session isn't working.	Return to the briefing. The session is only as deep as the quality of the WHERE YOU ARE STUCK field. A vague problem produces vague questions.

THE OVERRIDE LINE

"Do not rewrite the line. Ask me what my character is afraid to say."

THE DEPTH LINE

"Stay on this. Ask me again — differently."

Part Four

Most writers keep excavating past the substrate because they don't recognize it when it arrives. These four signals are distinct from surface and intermediate answers. When you observe two or more of them in the same response, stop. You are there.

01 The answer names a fear, not a trait

Surface and intermediate answers describe who the character is. The substrate names what they are protecting. "She is private" is a trait. "She is afraid that asking for help will prove she was never as strong as she told herself she was" is a substrate. The shift from description to fear is the signal.

02 The answer surprises you

You did not plan to write this sentence. It arrived. This is the difference between a constructed answer (generated by the part of you that already knew what it wanted to say) and a retrieved answer (surfaced by the part of you that has been watching your character without your conscious awareness). The surprise is the tell.

03 The answer accounts for behavior across the whole story

A substrate is a structural truth — it should explain not just the current scene but a pattern across the novel. If the sentence you wrote only explains one moment, you may have found a pressure point rather than the root. If it retroactively makes sense of multiple scenes and choices, you are at the substrate.

04 The answer slows you down

When you hit the substrate, your typing speed often drops. You may re-read the AI's question. You may pause before writing. This is not confusion — it is the quality of

attention that arrives when something real surfaces. Glib answers are produced quickly. Substrate answers require something from you.

Part Five

The session ends when you have the substrate sentence. What you do in the next thirty minutes is where the discovery earns its value.

Write the substrate sentence in your character notes. Word it as a fear, not a character description. This is a reference document — you will return to it every time you write a scene involving this character.

Do not rewrite the scene immediately. Read the original flat dialogue again — now, with the substrate visible. You are not looking for what to change yet. You are looking for what the character is not allowing themselves to say. Let the gap between the two become sharp before you touch the line.

Rewrite from the inside. The character should not say the substrate. They should say something that manages it — deflects it, circles it, protects it. The substrate is the thing beneath the words, not the words themselves. If the rewrite names the fear, you have overwritten. Back up.

Read the rewrite aloud. The test is not whether it sounds better. The test is whether the character now sounds like someone who has something to protect. The reader will not be able to name what that thing is. They will feel it.

A NOTE ON THE LONG GAME

A substrate sentence is a permanent asset. Once you have excavated a character's core fear, that sentence applies to every scene they appear in, for the entire novel. You do not re-excavate the same substrate. You use it. Return to the session only when you believe a scene has revealed a secondary substrate — a wound beneath a wound, which becomes relevant in later acts.

Part Six

Writers often believe that a session fails because of a poor question. Usually, it fails because of a poor briefing.

The AI's questions are generated from what you give it in the character briefing. A vague briefing produces generic questions. Generic questions produce generic answers. The session feels like a writing exercise rather than an excavation.

The three-field briefing from your 4.1 Training Guide is not a formality. The WHERE YOU ARE STUCK field in particular is the highest-leverage element in the entire session setup. Here is the distinction in practice:

WEAK BRIEFING (produces surface questions)	STRONG BRIEFING (produces depth questions)
"The dialogue feels flat and I can't figure out why."	"The character says she doesn't care if her son moves away. I know she does care. I don't know what she's actually afraid of."
"My character seems one-dimensional in this scene."	"He gives a perfectly logical answer when his colleague challenges him. The logic is true. But it doesn't feel like him — it feels managed. What is he not letting himself say?"
"This scene doesn't work."	"She apologizes in this scene and I don't believe it. I've written three versions. Something underneath the apology is wrong and I can't name it."

Before every session: read the WHERE YOU ARE STUCK field you've written and ask yourself — does this name the symptom, or does it name the gap between what I know and what I can access? If it only names the symptom, rewrite it before you open the AI.

RESEARCH GROUNDING

Cognitive scientists studying creative insight (Schooler & Melcher, 1995; Christoff et al., 2009) have demonstrated that framing a problem with specificity — naming what is known alongside what is unknown — activates broader associative retrieval than vague problem statements. The quality of the question you bring determines the depth of what surfaces.

QUICK REFERENCE

Session Setup Checklist

- ☐ Write your substrate goal: what fear are you trying to name?
- ☐ Write your three-field briefing — make WHERE YOU ARE STUCK specific, not symptomatic.
- ☐ Open a fresh AI session (don't continue a prior session).
- ☐ Paste your 4.1 coaching setup language, then paste the briefing immediately below it.
- ☐ Read the first question. Take time before you answer.
- ☐ If your first answer is quick and complete, flag it as surface — say so to the AI.
- ☐ If the AI rewrites or generates, use the override line and return it to questioning.
- ☐ When an answer surprises you or slows you down, write it in full and stop.
- ☐ Verify it names a fear, not a trait.
- ☐ Close the session and return to the dialogue.