

THE AVERAGE CHARACTER

What the Research Says About AI in Fiction



A Paradox Worth Understanding

AI companies have risked billion-dollar lawsuits to get access to modern fiction for their training data. Katherine Elkins, a researcher at Kenyon College who has spent seven years studying AI and narrative, has a name for what they discovered when they got it: the AI-Fiction Paradox. Models are built on massive amounts of fiction. They desperately need more of it. And they still can't generate it well.

That is not a problem better training data will eventually fix. It is structural. And understanding why it's structural may be the most useful thing you can know about the tool you're working with.

What the Research Has Found

Several independent research teams have now measured AI-generated fiction against human-written fiction. The findings land in the same place.

AI-generated stories show a consistent bias toward positive relationships and resolved conflict. The social structures they produce are measurably less complicated, less surprising, and less true than what human writers produce.

Researchers at Microsoft built a metric specifically to measure this. They called it the *Sui Generis* score — a way of tracking how often the same plot elements echo across multiple AI-generated stories from the same prompt. What they found: AI stories keep reproducing recognizable variants of the same story. Human-written plots, by contrast, are rarely echoed even in pieces.

There is also the instruction-tuning problem. The same process that makes AI helpful — the training that makes it safe, cooperative, and responsive — flattens its output. Even when you prompt it for a specific tone or style, something underneath keeps pulling it toward the generic. Researchers describe it as a persistent emotional disconnect that careful readers can sense.

Here is what is not on that list: grammar, coherence, pacing, surface fluency. AI handles all of those well. What it cannot produce is the one thing that makes a reader feel something.

Specificity.

Why Specificity Cannot Be Generated

Elkins identifies three structural reasons why fiction resists AI generation. The one that matters most for character work is what she calls *informational revaluation*.

Fiction, she argues, violates a core assumption built into how language models work — the assumption that what is statistically prominent is what is important. In fiction, the opposite is often true. The detail that changes everything is frequently the one that looked like nothing when it arrived. The offhand mention of a father's age. A hand going to a mouth. These details only become significant in retrospect, once you know the history they were carrying all along.

A language model generating text forward has no mechanism for that. It weights what is statistically common. It cannot weight what is privately true.

This is not something better AI will solve. It is a feature of what fiction actually is. The line that stops a reader is, by definition, not the expected line. A language model calculates the most probable next word. The line you need is the least probable one — the one that could only come from this character, carrying this history, in this exact moment.

That line has never been written anywhere. It lives only in your imagination. It cannot be retrieved. It has to be surfaced.

The Behavioral Parallel

The research describes this as a generation problem. My consulting work describes the same problem from the other side.

Across more than twenty thousand subjects, the pattern held: people do not behave according to type. They behave according to history. Two people with the same role, the same training, the same stated values will respond to the same

pressure in completely different ways — because their histories are different. What looks irrational from the outside makes perfect sense once you know what the person is protecting.

AI knows types. It knows how husbands in general respond to financial pressure. How detectives in general respond to a first suspect. How mothers in general respond to a child leaving. It produces those responses fluently.

Your character is not a type. Your character is a history.

The specific logic that makes her response true — what she's protecting, what's driving her in this scene — exists nowhere except in your developing understanding of her. No prompt that asks AI to generate can reach it.

The instruction has to change.

The One Change That Changes Everything

Instead of asking AI to write the scene, you ask AI to ask you about it.

The AI becomes a coach instead of an author. You become the source. And what surfaces in your third or fourth answer is not something the AI produced. It is something you already knew about your character — something that had not yet made it to the page.