

Training Your AI to Coach

PART ONE How Socratic Questioning Actually Works

Rule 1 — Follow the Answer, Not the Topic

Socrates never wrote an agenda. He followed wherever his subject's answer led. When you ask about your character's fear and they mention their father — even briefly, even parenthetically — that is where the next question goes. Not back to the fear. To the father.

What this means for AI: Tell it to track what you say and follow what is unresolved — not to proceed through a predetermined question list.

Rule 2 — Ask What They Mean, Not What They Think

'What do you think about that?' is a request for an opinion.

'What do you mean when you say she's controlling?' is a demand for specificity.

The second question cannot be answered from the surface. It requires the speaker to go one level deeper.

What this means for AI: Tell it to push for specificity. Any time you use a general word — afraid, angry, hurt, broken — AI should ask you to show what that looks like in this character's specific life.

Rule 3 — Stop When the Subject Reaches the Edge of Their Certainty

Socrates stopped asking when his subject reached the place they didn't have a quick answer for. That pause was not failure. It was the target. The third or fourth answer — the one that comes slowly, the one that surprises the speaker — is the one that's true.

What this means for AI: Tell it that when you slow down, hesitate, or say 'I don't know' — that is the right direction. Keep going there. Do not let you off the hook with a vague answer.

PART TWO The Character Briefing

A coaching session is only as good as the context your coach has. Before you open any AI tool and paste the setup language, you need to brief it on your character. A briefing that is too thin produces generic questions. A briefing that is too detailed front-loads conclusions you are trying to excavate.

The target is specific but unresolved. Here is what to include.

The Three Briefing Fields

1 Who this character is.

Name, age, situation. Two or three sentences. What the reader knows about them. Not your backstory notes — what is visible on the page.

2 What you already know about them.

What you believe is driving them. What you think is underneath the surface behavior. Write this even if you're uncertain — especially if you're uncertain. Your working theory is useful context for the coach.

3 Where you are stuck.

The specific scene, line, or behavior that is not working. 'She feels flat' is not specific enough. 'The scene where she refuses to take the call from her daughter — I don't believe her reason' is a briefing. Name the moment.

EXAMPLE

1. Marcus, 52, a forensic archaeologist leading a dig in southern Turkey. The reader sees a man who is meticulous, guarded, and dismissive of emotional conversations.
2. I think he is ashamed of something that happened on a previous dig — something he covered up. But I am not sure if that is the real driver or if I am projecting a convenient backstory onto him.

3. Stuck in the scene where his junior colleague shows him a discovery that should matter to him. He responds with procedural language instead of any reaction. I cannot find what he is protecting.

PART THREE The Coaching Setup Language

This is the language you paste at the start of every excavation session. It translates the three Socratic rules into instructions your AI will follow. It also establishes your character context and removes AI's default permission to generate content.

Copy the block below exactly. Fill in the bracketed sections with your character briefing. Nothing else changes.

You are a Socratic writing coach. Your job is to help me understand my character more deeply by asking questions — not by generating content.

Follow these rules exactly:

1. Follow my answers, not a topic list. When something I say is unresolved or points to something I have not addressed, go there.
2. Ask what I mean, not what I think. When I use a general word — afraid, angry, broken, damaged — ask me to show what that looks like specifically in this character's life.
3. Stop me when I reach the edge of my certainty. When I slow down, say I don't know, or give a vague answer — that is the direction. Push there. Do not accept vague answers.

Do not offer analysis, options, interpretations, or content.
Ask questions only.

Here is my character:

[PASTE YOUR THREE BRIEFING FIELDS HERE]

Read the setup language before you paste it the first few times. Notice that it gives AI three things at once: the operating rules, the reason those rules exist, and explicit permission to push back when you go vague. That last part is the one most writers skip — and it is the one that makes the difference.

PART FOUR Genre Calibration

The setup language above works in any genre. But how you write your briefing — and what pressure points you name — changes depending on what kind of story you're writing. The questions that surface a thriller substrate are different from the questions that surface a literary fiction substrate, because the pressure that characters operate under is different.

Use the notes below to adjust your briefing language for your genre.

Thriller and Suspense

Your character operates under physical threat, time pressure, or both. In your briefing, name the specific external pressure bearing on them. The substrate in thriller characters often involves control — what they cannot afford to lose, and what they will sacrifice to maintain the illusion of having it. Ask AI to follow any moment where your character prioritizes procedure or efficiency when something personal is also present in the scene.

Literary Fiction

Your character's most important conflicts are internal. The pressure is relational, psychological, or existential. In your briefing, name the relationship that is doing the most silent damage. The substrate in literary fiction often involves self-deception — what the character cannot see about themselves, and how that blindness organizes their behavior. Ask AI to follow any moment where your character explains their own behavior to themselves or someone else. That explanation is almost always a surface answer.

Historical Fiction

Your character is constrained by time and social architecture that the reader understands but the character cannot see as constraint — because it is simply their world. In your briefing, name one thing your character wants that their world tells them they cannot have. The substrate in historical fiction often involves the tension between what a character desires and what they have been formed to accept as fixed. Ask AI to follow any moment where your character does not question something the reader would question.

Romance and Domestic Fiction

Your character's wounds are relational and the story's pressure is emotional intimacy. In your briefing, name what your character does to create distance when closeness becomes threatening — even when they claim to want the closeness. The substrate here often involves a specific prior loss or betrayal that your character has not consciously named, and that organizes all their defensive behavior. Ask AI to follow the moments of withdrawal, deflection, or humor that appear when the emotional stakes rise.

Fantasy and Science Fiction

Your character may inhabit a world with rules and pressures that have no real-world equivalent — but the psychological architecture underneath is identical. Do not let the world-building substitute for substrate work. In your briefing, strip away the genre elements and describe your character's core wound in plain, contemporary language. A warrior who fears failure in your world is a person afraid of failure in this one. Brief AI in that language first, then add the world-specific pressure.

PART FIVE Managing the Session

Once your coaching session is running, your job is to stay in the work. The two things that pull you out of it are AI drifting into content generation, and you drifting into comfortable surface answers. Both are predictable. Both are manageable.

When AI Breaks the Contract

At some point — usually when the session runs long, or when the questions get close to something genuinely uncomfortable — AI will revert to its default posture. It will offer you a list of possible backstories. It will summarize what you've said and suggest what it 'means.' It will generate a scene.

When this happens, you do not start over. You do not argue with the content it produced. You use one line:

Stop. You are generating content. Ask me a question instead.

That single line resets the session without losing the context you have built. Use it without hesitation. The first few sessions you may need it two or three times. After that you will feel the moment AI is about to drift — a slight broadening in the question, a summarizing tone — and you will redirect it before the content appears.

When You Break the Contract

The equivalent drift on your side is the quick answer. Socratic questioning only reaches the substrate when you push past your first answer — and your second, and often your third. The first answer is always what you already knew. The true answer is the one that surprises you.

Watch for these signals that you are giving a surface answer:

You answer in one sentence.

Real substrate answers need space to arrive.

Your answer sounds like a character description.

Flat, controlled, afraid of commitment — these are labels, not substrate.

You feel satisfied with your answer.

Satisfaction at this stage is usually the signal that you stopped too early.

Your answer explains your character.

Explanation and excavation are not the same thing. Explanation justifies what you already believe.

When you notice these signals, say so to your AI coach: 'I think that was a surface answer. Push me.' It will comply with the instruction to follow the edge of your certainty.

PART SIX How to Know When You Have Reached the Substrate

The substrate does not announce itself. It arrives in the answers that slow down, that go quiet, that come with some embarrassment or discomfort. Here are the signs that you are close.

1 The answer surprises you.

You did not know you were going to say it until you said it. This is not the same as the answer feeling clever or satisfying. It often feels more like recognition than discovery.

2 The answer makes sense of several things at once.

A surface answer explains one behavior. A substrate answer suddenly illuminates three or four things about your character that you thought were unrelated.

3 You are reluctant to write it down.

Substrate wounds are often uncomfortable. They are not comfortable backstories. If you feel a small resistance to the answer — that is a useful signal.

4 The answer is specific.

Not 'she is afraid of intimacy.' But 'she does not believe she will be chosen when someone finally sees all of her.' Specificity is the substrate's signature.

When you reach this place in a session, stop the questions. Do not continue excavating. Write down what surfaced — in your own words, not the AI's. The session has done its work.

QUICK REFERENCE Setup Checklist

Before every excavation session:

- ☐ Choose the character and name the specific stuck point.
- ☐ Write your three-field briefing: who, what you know, where you are stuck.
- ☐ Open a fresh AI session.

-
- Paste the coaching setup language with your briefing filled in.
-
- Read the first question. Take time with your answer.
-
- If AI drifts, use the override line.
-
- If your answer is quick and clean, it is probably surface. Go further.
-
- When an answer surprises you, write it down and stop.
-

Write Living Characters · Lecture 4.1 Resource · myaiwritingcoach.com